

MCP Phase 19-20 Context and Decisions Summary

Purpose

This document summarizes the key conclusions, decisions, and guidance from the current chat regarding ChatGPT project architecture, NAS file organization, and RAG/DB strategy. It is intended as a handoff reference for continuation of Phase 19 (Hardening) and Phase 20 (NAS Database Integration) work within the **MCP for Spear Enterprise LLC - Tier 1** project.

1. ChatGPT Project Architecture (Context)

- ChatGPT **Projects are flat containers**. There is **no native support** for:
 - Subfolders
 - Nested projects
 - Parent/child project hierarchies
- A “Portfolio Project” is **not a special feature**. It is simply:
 - A newly created project
 - Run with stricter governance rules (index, decisions, summaries only)
 - Optionally configured with **project-only memory** (must be set at creation time or via sharing)
- Because nesting is unavailable and no public OpenAI roadmap exists for nested projects, **external structure must be enforced outside ChatGPT**.

Conclusion: ChatGPT Projects are UI workspaces only. The real hierarchy and memory discipline must live on the NAS and in the Phase 20 database.

2. Portfolio Pattern (Operating Model)

The adopted model is a **Portfolio Pattern**, which simulates nesting via discipline:

- One top-level **Portfolio / Boardroom project**
- Index, decision log, risk register only
- Separate **Program Hub projects** (Heber Campus, MCP, OSY, etc.)
- Separate **Workstream projects** for active iteration

Critical rule: - No raw thinking in the Portfolio project - Only promoted, decision-grade outputs move upward

This pattern remains valid even if ChatGPT later adds nesting.

3. Tier 1 NAS Structure (Authoritative Source)

Your Tier 1 NAS structure defines `/SE_T1/` as the **single immutable root of truth**.

Key points from governance documents: - MCP is the **infrastructure backbone**, not a content project - ATB-approved artifacts are mirrored into `/SE_T1/GOVERNANCE/` - Lifecycle separation is expected (WORKING / FINAL / ARCHIVE)

Implication: - File organization and truth ownership are enforced at the NAS level - ChatGPT projects reference this structure, not replace it

4. Phase 19 Hardening – Required Controls

Phase 19 must lock governance before Phase 20 ingestion.

Required by end of Phase 19:

1. **Lifecycle Lanes (Non-Negotiable)**
2. `SOURCE` – raw imports, vendor PDFs, dumps
3. `WORKING` – drafts, in-flight analysis
4. `FINAL` – decision-grade, authoritative
5. `ARCHIVE` – superseded finals

6. **Retrieval Priority Rule (for Agents / RAG)**
7. FINAL first
8. WORKING only if FINAL has no hits
9. SOURCE as evidence only, not narrative truth

10. **Naming & Revision Enforcement (on promotion only)**
11. ASCII only
12. Explicit REV-A / REV-B
13. Do not rename legacy files during import

14. **Continuity Registry Discipline**

-
- 15. All ingest, promotion, and supersession events append to registry/log
-

5. Phase 20 – NAS Database Strategy

Purpose of the DB

The Phase 20 DB is the **real hierarchy** and **anti-fragmentation layer**: - Stable IDs - Version control - Relationship mapping - Deterministic retrieval for agents

Folders are for humans. The DB is for machines.

DB Placement (per Tier 1 governance)

Place DB infrastructure under MCP:

- /SE_T1/INFRASTRUCTURE_CORE/MCP/DB/
- index/ (SQLite initially)
- records/ (JSON metadata per artifact)
- objects/ (content-addressed storage by hash, if used)
- ingest/ (atomic drop zone)
- Continuity and audit logs under:
 - /SE_T1/INFRASTRUCTURE_CORE/CONTINUITY_REGISTRY/

Content remains in its owning branches (SDC_COMMS, ENERGY_DIVISION, etc.).

6. File Migration & Placement Strategy (500 Files)

Key Decision

Do not reorganize first. Mirror first.

Recommended Steps:

1. Mirror-copy existing SDC-COMMS content into:
 - 2. /SE_T1/SPECIAL_PROJECTS/SDC_COMMS/SOURCE/RAW_IMPORT/<YYYYMMDD>/...
3. Generate a manifest:
4. Full path
5. File type

6. Modified time
 7. Size
 8. SHA-256 hash
 9. Deduplicate by hash
 10. One canonical candidate per duplicate group
 11. Classify by lifecycle (SOURCE / WORKING / FINAL)
 12. Stand up DB records pointing to existing paths (no renames)
 13. Promote only top-value authoritative docs (30-50) into FINAL/
-

7. Canonical Artifact Model (Anti-RAG Fragmentation)

Each artifact must have a **stable ID**, independent of filename or folder.

Minimum DB fields: - artifact_id - title - owning_branch - status (source | working | final | archived | governance) - revision - sha256 - canonical_path - source_path - supersedes - tags

Hard rule: - Only one current FINAL per artifact_id

This rule alone eliminates most RAG hallucination and contradiction issues.

8. Domain-Based Organization (Not Project-Based)

For PUBLISHED/FINAL content, organize by **domain**, not by project:

- GOVERNANCE
- ARCHITECTURE
- SECURITY
- OPERATIONS
- ENGINEERING (POWER, H2, DATA-CENTER)
- FINANCE
- LEGAL

Projects are expressed via DB tags, not folder depth.

9. Bottom-Line Guidance for Next Chat

- Phase 19 is about **locking rules**, not moving files
 - Phase 20 DB is the **real hierarchy and memory system**
 - RAG quality depends on lifecycle separation, not chunking tricks
 - Only FINAL + GOVERNANCE should be default retrieval truth
 - The Tier 1 NAS structure already supports this—do not fight it
-

Open Items for Next Chat

- Finalize Phase 19 hardening checklist (explicit gates)
 - Define Phase 20 DB v0 schema (SQLite)
 - Decide initial promotion set (top 30–50 artifacts)
 - Define agent retrieval policy in MCP (scope + ranking)
-

End of Summary